

# Python Lab Exercise #2

## Objectives:

- Load .csv files into `pandas` `DataFrames`
- Describe and manipulate data in `Series` and `DataFrames`
- Visualize data using `DataFrame` methods and `matplotlib`



In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

## What is Pandas?

Pandas, as [the Anaconda docs](#) tell us, offers us "High-performance, easy-to-use data structures and data analysis tools." It's something like "Excel for Python", but it's quite a bit more powerful.

Let's read in the heart dataset.

Pandas has many methods for reading different types of files. Note that here we have a .csv file.

Read about this dataset [here](#).

In [2]:

```
heart_df = pd.read_csv('data/heart.csv')
```

The output of the `.read_csv()` function is a pandas *DataFrame*, which has a familiar tabular structure of rows and columns.

In [3]:

```
type(heart_df)
```

Out[3]:

```
pandas.core.frame.DataFrame
```

In [4]:

```
heart_df
```

Out[4]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
298	57	0	0	140	241	0	1	123	1	0.2	1	0	3	0
299	45	1	3	110	264	0	1	132	0	1.2	1	0	3	0
300	68	1	0	144	193	1	1	141	0	3.4	1	2	3	0
301	57	1	0	130	131	0	1	115	1	1.2	1	1	3	0
302	57	0	1	130	236	0	0	174	0	0.0	1	1	2	0

303 rows × 14 columns

## DataFrames and Series

Two main types of pandas objects are the **DataFrame** and the **Series**, the latter being in effect a single column of the former:

In [5]:

```
age_series = heart_df['age']
type(age_series)
```

Out[5]:

```
pandas.core.series.Series
```

Notice how we can isolate a column of our **DataFrame** simply by using square brackets together with the name of the column.

Both **Series** and **DataFrames** have an *index* as well:

In [6]:

```
heart_df.index
```

Out[6]:

```
RangeIndex(start=0, stop=303, step=1)
```

In [7]:

```
age_series.index
```

Out[7]:

```
RangeIndex(start=0, stop=303, step=1)
```

**Pandas is built on top of NumPy**, and we can always access the **NumPy** array underlying a **DataFrame** using `.values`.

In [8]:

```
heart_df.values
```

```
Out[8]:
array([[63., 1., 3., ..., 0., 1., 1.],
       [37., 1., 2., ..., 0., 2., 1.],
       [41., 0., 1., ..., 0., 2., 1.],
       ...,
       [68., 1., 0., ..., 2., 3., 0.],
       [57., 1., 0., ..., 1., 3., 0.],
       [57., 0., 1., ..., 1., 2., 0.]])
```

Basic DataFrame Attributes and Methods

```
.head()
```

```
In [9]:
```

```
heart_df.head()
```

Out[9]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1

```
.tail()
```

```
In [10]:
```

```
heart_df.tail()
```

Out[10]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
298	57	0	0	140	241	0	1	123	1	0.2	1	0	3	0
299	45	1	3	110	264	0	1	132	0	1.2	1	0	3	0
300	68	1	0	144	193	1	1	141	0	3.4	1	2	3	0
301	57	1	0	130	131	0	1	115	1	1.2	1	1	3	0
302	57	0	1	130	236	0	0	174	0	0.0	1	1	2	0

```
.info()
```

```
In [11]:
```

```
heart_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
1   sex         303 non-null   int64
2   cp          303 non-null   int64
```

```
3    trestbps    303 non-null    int64
4    chol        303 non-null    int64
5    fbs         303 non-null    int64
6    restecg     303 non-null    int64
7    thalach     303 non-null    int64
8    exang       303 non-null    int64
9    oldpeak     303 non-null    float64
10   slope       303 non-null    int64
11   ca          303 non-null    int64
12   thal        303 non-null    int64
13   target      303 non-null    int64
```

```
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

`.describe()`

In [12]:

```
heart_df.describe()
```

Out[12]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	0.528053	149.646865	0.326733	1.0396
std	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	0.525860	22.905161	0.469794	1.1610
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000	71.000000	0.000000	0.0000
25%	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	0.000000	133.500000	0.000000	0.0000
50%	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.000000	153.000000	0.000000	0.8000
75%	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000	1.000000	166.000000	1.000000	1.6000
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000	202.000000	1.000000	6.2000

`.dtypes`

In [13]:

```
heart_df.dtypes
```

Out[13]:

```
age          int64
sex          int64
cp           int64
trestbps     int64
chol         int64
fbs          int64
restecg      int64
thalach      int64
exang        int64
oldpeak      float64
slope        int64
ca           int64
thal         int64
target       int64
dtype: object
```

`.shape`

In [14]:

```
heart_df.shape
```

Out[14]:


(303, 14)

## Exploratory Plots

Let's make ourselves a histogram of ages:

In [49]:


```
sns.set_style('darkgrid')
#sns.distplot(a=heart_df['age']);
# For more recent versions of seaborn:
sns.histplot(data=heart_df['age'], kde=True);
```



And while we're at it let's do a scatter plot of maximum heart rate vs. age:

In [16]:

```
sns.scatterplot(x=heart_df['age'], y=heart_df['thalach']);
```



## Adding to a DataFrame

### Adding Rows

Here are two rows that our engineer accidentally left out of the .csv file, expressed as a Python dictionary:

In [17]:

```
extra_rows = {'age': [40, 30], 'sex': [1, 0], 'cp': [0, 0], 'trestbps': [120, 130],
              'chol': [240, 200],
              'fbs': [0, 0], 'restecg': [1, 0], 'thalach': [120, 122], 'exang': [0, 1],
```

```
'oldpeak': [0.1, 1.0], 'slope': [1, 1], 'ca': [0, 1], 'thal': [2, 3],  
'target': [0, 0]}
```

```
extra_rows
```

```
Out[17]:
```

```
{'age': [40, 30],  
 'sex': [1, 0],  
 'cp': [0, 0],  
 'trestbps': [120, 130],  
 'chol': [240, 200],  
 'fbs': [0, 0],  
 'restecg': [1, 0],  
 'thalach': [120, 122],  
 'exang': [0, 1],  
 'oldpeak': [0.1, 1.0],  
 'slope': [1, 1],  
 'ca': [0, 1],  
 'thal': [2, 3],  
 'target': [0, 0]}
```

How can we add this to the bottom of our dataset?

```
In [18]:
```

```
# Let's first turn this into a DataFrame.  
# We can use the .from_dict() method.
```

```
missing = pd.DataFrame(extra_rows)  
missing
```

```
Out[18]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	40	1	0	120	240	0	1	120	0	0.1	1	0	2	0
1	30	0	0	130	200	0	0	122	1	1.0	1	1	3	0

```
In [19]:
```

```
# Now we just need to concatenate the two DataFrames together.  
# Note the `ignore_index` parameter! We'll set that to True.
```

```
heart_augmented = pd.concat([heart_df, missing],  
                             ignore_index=True)
```

```
In [20]:
```

```
# Let's check the end to make sure we were successful!
```

```
heart_augmented.tail()
```

```
Out[20]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
300	68	1	0	144	193	1	1	141	0	3.4	1	2	3	0
301	57	1	0	130	131	0	1	115	1	1.2	1	1	3	0
302	57	0	1	130	236	0	0	174	0	0.0	1	1	2	0
303	40	1	0	120	240	0	1	120	0	0.1	1	0	2	0
304	30	0	0	130	200	0	0	122	1	1.0	1	1	3	0

## Adding Columns

Adding a column is very easy in `pandas`. Let's add a new column to our dataset called "test", and set all of its values to 0.

values to 0.

In [21]:

```
heart_augmented['test'] = 0
```

In [22]:

```
heart_augmented.head()
```

Out[22]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target	test
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1	0
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1	0
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1	0
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1	0
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1	0

I can also add columns whose values are functions of existing columns.

Suppose I want to add the cholesterol column ("chol") to the resting systolic blood pressure column ("trestbps"):

In [23]:

```
heart_augmented['chol+trestbps'] = heart_augmented['chol'] + heart_augmented['trestbps']
```

In [24]:

```
heart_augmented.head()
```

Out[24]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target	test	chol+trestbps
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1	0	378
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1	0	380
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1	0	334
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1	0	356
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1	0	474

## Filtering

We can use filtering techniques to see only certain rows of our data. If we wanted to see only the rows for patients 70 years of age or older, we can simply type:

In [25]:

```
heart_augmented['age'] >= 70
```

Out[25]:

```
0      False
1      False
2      False
3      False
4      False
...
300     False
301     False
302     False
303     False
304     False
```

```
304         raise
Name: age, Length: 305, dtype: bool
```

In [26]:

```
heart_augmented[heart_augmented['age'] >= 70]
```

Out[26]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target	test	chol+trestbps
25	71	0	1	160	302	0	1	162	0	0.4	2	2	2	1	0	462
60	71	0	2	110	265	1	0	130	0	0.0	2	1	2	1	0	375
129	74	0	1	120	269	0	0	121	1	0.2	2	1	2	1	0	389
144	76	0	2	140	197	0	2	116	0	1.1	1	0	2	1	0	337
145	70	1	1	156	245	0	0	143	0	0.0	2	0	2	1	0	401
151	71	0	0	112	149	0	1	125	0	1.6	1	0	2	1	0	261
225	70	1	0	145	174	0	1	125	1	2.6	0	0	3	0	0	319
234	70	1	0	130	322	0	0	109	0	2.4	1	3	2	0	0	452
238	77	1	0	125	304	0	0	162	1	0.0	2	3	2	0	0	429
240	70	1	2	160	269	0	1	112	1	2.9	1	1	3	0	0	429

Use '&' for "and" and '|' for "or".

## Exercise

Display the patients who are 70 or over as well as the patients whose trestbps score is greater than 170.

In [27]:

```
# Enter your code here
age_trestbps = heart_df[(heart_df['age'] >= 70) | (heart_df['trestbps'] > 170)]
age_trestbps
```

Out[27]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target	
	8	52	1	2	172	199	1	1	162	0	0.5	2	0	3	1
	25	71	0	1	160	302	0	1	162	0	0.4	2	2	2	1
	60	71	0	2	110	265	1	0	130	0	0.0	2	1	2	1
	101	59	1	3	178	270	0	0	145	0	4.2	0	0	3	1
	110	64	0	0	180	325	0	1	154	1	0.0	2	0	2	1
	129	74	0	1	120	269	0	0	121	1	0.2	2	1	2	1
	144	76	0	2	140	197	0	2	116	0	1.1	1	0	2	1
	145	70	1	1	156	245	0	0	143	0	0.0	2	0	2	1
	151	71	0	0	112	149	0	1	125	0	1.6	1	0	2	1
	203	68	1	2	180	274	1	0	150	1	1.6	1	0	3	0
	223	56	0	0	200	288	1	0	133	1	4.0	0	2	3	0
	225	70	1	0	145	174	0	1	125	1	2.6	0	0	3	0
	234	70	1	0	130	322	0	0	109	0	2.4	1	3	2	0
	238	77	1	0	125	304	0	0	162	1	0.0	2	3	2	0
	240	70	1	2	160	269	0	1	112	1	2.9	1	1	3	0
	241	59	0	0	174	249	0	1	143	1	0.0	1	0	2	0




	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
248	54	1	1	192	283	0	0	195	0	0.0	2	1	3	0
260	66	0	0	178	228	1	1	165	1	1.0	1	2	3	0
266	55	0	0	180	327	0	2	117	1	3.4	1	0	2	0

## Exploratory Plot

Using the subframe we just made, let's make a scatter plot of their cholesterol levels vs. age and color by sex:

In [28]:

```
at_risk = age_trestbps
sns.scatterplot(data=at_risk, x='age', y='chol', hue='sex');
```



## .loc and .iloc

We can use `.loc` to get, say, the first ten values of the age and resting blood pressure ("trestbps") columns:

In [29]:

```
heart_augmented.loc
```

Out[29]:  
<pandas.core.indexing.\_LocIndexer at 0x7fbdb1544180>

In [30]:

```
heart_augmented.loc[:9, ['age', 'trestbps']]
```

Out[30]:

	age	trestbps
0	63	145
1	37	130
2	41	130
3	56	120
4	57	120
5	57	140
6	56	140
7	44	120
8	52	172
9	57	150

`.iloc` is used for selecting locations in the DataFrame by number:

In [31]:

```
heart_augmented.iloc
```

Out[31]:

```
<pandas.core.indexing._iLocIndexer at 0x7fbdb153e400>
```

In [32]:

```
heart_augmented.iloc[3, 0]
```

Out[32]:

56

In [33]:

```
heart_augmented.head()
```

Out[33]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target	test	chol+trestbps
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1	0	378
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1	0	380
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1	0	334
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1	0	356
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1	0	474

## Exercise

How would we get the same slice as just above by using `.iloc()` instead of `.loc()`?

In [34]:

```
# Enter your code here
heart_augmented.iloc[:5, 0:]
```

Out[34]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target	test	chol+trestbps
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1	0	378
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1	0	380
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1	0	334
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1	0	356
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1	0	474

## Statistics

`.mean()`

In [35]:

```
heart_augmented.mean()
```

Out[35]:

age 54.239344

```

age                0.681967
sex                0.960656
cp                131.580328
trestbps          246.091803
chol              0.147541
fbs               0.527869
restecg           149.459016
thalach           0.327869
exang             1.036393
oldpeak           1.396721
slope             0.727869
ca               2.314754
thal             0.540984
target            0.000000
test              377.672131
chol+trestbps
dtype: float64

```

**Be careful! Some of these will are not straightforwardly interpretable. What does an average "sex" of 0.682 mean?**

`.min()`

In [36]:

```
heart_augmented.min()
```

Out[36]:

```

age                29.0
sex                0.0
cp                0.0
trestbps          94.0
chol              126.0
fbs               0.0
restecg           0.0
thalach           71.0
exang             0.0
oldpeak           0.0
slope             0.0
ca               0.0
thal             0.0
target            0.0
test              0.0
chol+trestbps     249.0
dtype: float64

```

`.max()`

In [37]:

```
heart_augmented.max()
```

Out[37]:

```

age                77.0
sex                1.0
cp                3.0
trestbps          200.0
chol              564.0
fbs               1.0
restecg           2.0
thalach           202.0
exang             1.0
oldpeak           6.2
slope             2.0
ca               4.0
thal             3.0
target            1.0
test              0.0

```

```
chol+trestbps      679.0  
dtype: float64
```

## Series Methods

`.value_counts()`

How many different values does slope have? What about sex? And target?

In [38]:

```
heart_augmented['slope'].value_counts()
```

Out[38]:

```
2    142  
1    142  
0     21  
Name: slope, dtype: int64
```

In [39]:

```
heart_augmented['sex'].value_counts()
```

Out[39]:

```
1    208  
0     97  
Name: sex, dtype: int64
```

`.sort_values()`

In [40]:

```
heart_augmented['age'].sort_values()
```

Out[40]:


```
72     29  
304    30  
58     34  
125    34  
65     35  
..  
25     71  
60     71  
129    74  
144    76  
238    77  
Name: age, Length: 305, dtype: int64
```


## pandas -Native Plotting

The `.plot()` and `.hist()` methods available for DataFrames use a wrapper around `matplotlib`:

In [41]:

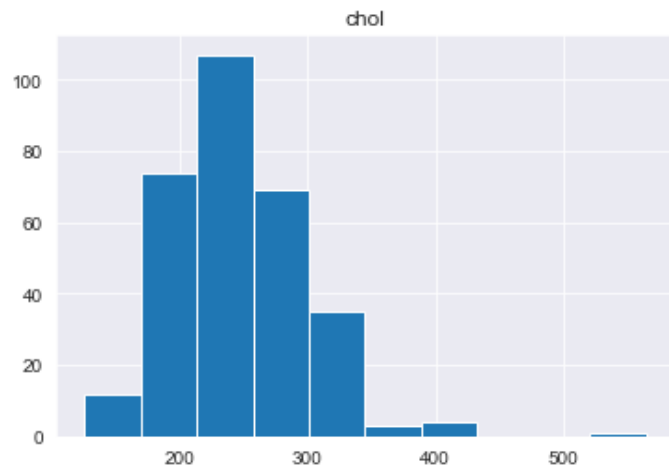
```
heart_augmented.plot(x='age', y='trestbps', kind='scatter');
```





In [42]:

```
heart_augmented.hist(column='chol');
```

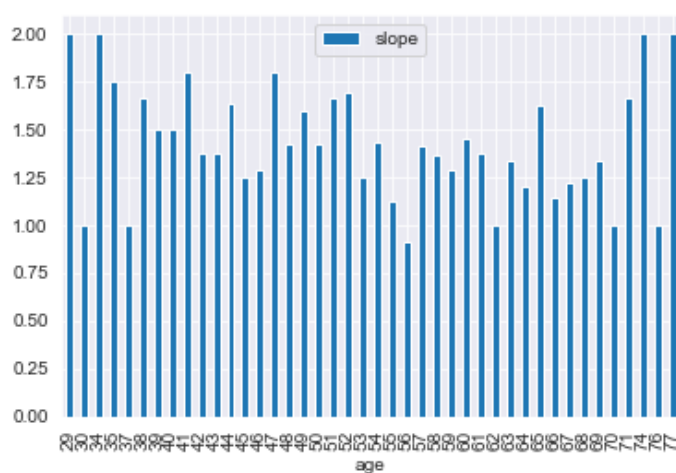


## Exercises

1. Make a bar plot of "age" vs. "slope" for the `heart_augmented` DataFrame.

In [43]:

```
# Enter your code here
new_heart_augmented = heart_augmented.groupby('age').mean()
new_heart_augmented.plot.bar(y='slope');
```




1. Make a histogram of ages for just the men in `heart_augmented` (`heart_augmented['sex']==1`).

In [44]:

```
# Enter your code here
men = heart_augmented[heart_augmented['sex'] == 1]
men.hist(column='age');
```





1. Make separate scatter plots of cholesterol vs. resting systolic blood pressure for the target=0 and the target=1 groups. Put both plots on the same figure and give each an appropriate title.

In [45]:

```
# Enter your code here
sns.scatterplot(data=heart_augmented, x='chol', y='trestbps', hue='target');
```

